

Multi Rater Modelling

Meritxell Riera i Marín, PhD candidate

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Motivation

Data Uncertainty

In Medical Imaging, there is inherent ambiguity in the annotations → aleatoric uncertainty

↓



Interrater Variability




Disagreement is common – ill-defined boundaries (e.g. Tumor margins)




↓

Exploit this variability

Depending on the disagreement, different approaches might be required







Motivation

Naively choosing one annotation (either raw or created from others) can bias training and evaluation.

Motivation

Naively choosing one annotation (either raw or created from others) can bias training and evaluation.

How do we merge, model or exploit multiple labels?

What is the impact on training and validation?

How does this affect calibration and reliability?

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Label Fusion – Overview

Naive Extremes

Simple
Consensus

Random
Sampling

Soft Labels

Probabilistic
Modeling
(STAPLE)

Other Methods

Label Fusion – Naive Extremes

- Union, “captures everything”
 - High sensitivity, low specificity
 - Oversegment
- Intersection, “only agree when sure”
 - Low sensitivity, high specificity
 - Undersegment




Easy to understand/implement



Oversimplify uncertainty → either too permissive or too conservative



Often used as baselines



Label Fusion – Simple Consensus

- Majority Voting (MV)
- Median = MV
 - Binary Cases
 - Odd number of raters




Robust to outliers



Still a hard label → ignores degree of uncertainty



Weighted voting → acknowledges rater expertise (e.g. juniors versus seniors)



Label Fusion – Random Sampling

- Instead of collapsing raters → sample one rater's label at random each epoch/iteration
- Preserves diversity across raters
- Forces model to “see” all annotation styles over training




Simple, no fusion required
Effective when disagreement is high



Adds variance → requires long training to
balance exposure to all raters



Often used as a regularizer → prevents overfitting
to a single GT



Label Fusion – Soft labels

- Soft Consensus: Instead of forcing a hard “truth,” use **soft probabilities** (fraction of raters agreeing).
- Can add **spatial smoothing** → uncertainty higher near boundaries.



Uncertainty-aware labels




Harder evaluation → soft compatible metrics




Closer to the **real distribution of expert uncertainty**, can help calibration

Individual Annotators




Soft-Consensus

SVLS






Spatially Varying
Label Smoothing


Label Fusion – STAPLE





Initialize: guess
true mask
(e.g. Majority
voting)



Label Fusion – STAPLE




Initialize: guess
true mask
(e.g. Majority
voting)






E-step: given sens/spec,
compute probability
each voxel is “true” →
probabilistic map.


M-step: given
current “truth,”
update sens/spec for
each rater.



Label Fusion – STAPLE




Initialize: guess true mask (e.g. Majority voting)



E-step: given sens/spec, compute probability each voxel is “true” → probabilistic map.

M-step: given current “truth,” update sens/spec for each rater.



When converges

Label Fusion – STAPLE

- Simultaneous Truth And Performance Level Estimation (STAPLE)
- Learns both a latent “true mask” and rater reliabilities




Accounts for **individual rater reliability**



Assume independence, voxel-wise independence may break for segmentation



Acknowledge raters are noisy observers



Label Fusion – Other Methods

- **DL-based label fusion:** CNNs that learn to combine rater maps into consensus
 - Small network learns to fuse multi-rater inputs into a consensus map
- **Multi-Annotator Training** – e.g. loss functions that compare predictions to *all raters* instead of a fused mask
- **Bayesian / Uncertainty-Aware Networks** – explicitly predict per-voxel uncertainty calibrated to rater disagreement.



Can incorporate spatial context & rater-specific patterns
Directly train model with uncertainty included



Require more data and careful design → harder to interpret and validate than classical methods



These are more flexible but less standardized → ongoing research

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Method wrt Disagreement

Method	Low disagreement (raters mostly agree)	High disagreement (raters split)
Union	$\approx$ Majority/Intersection $\rightarrow$ stable	Over-segmented, very sensitive
Intersection	$\approx$ Majority $\rightarrow$ stable	Under-segmented, overly strict
Majority Voting	Stable, near ground truth	Hides minority opinions $\rightarrow$ biased
Random Sampling	$\approx$ Majority $\rightarrow$ stable	Preserves diversity, robust
STAPLE	Works well $\rightarrow$ estimates reliabilities	Still OK, but limited by independence assumption
Soft Labels	Almost same as hard mask	Best at reflecting real uncertainty

When disagreement is low $\rightarrow$ all methods converge, choice doesn't matter much.
When disagreement is high $\rightarrow$ choice of fusion method **critically changes the "truth"** the model learns.

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Soft vs Hard on Training/Validation

Hard Labels		Soft Labels
Training	 • Faster convergence • Clear GT target • Simple loss functions (cross-entropy, Dice)	• Let the model learn uncertainty and variability • Reduces overconfidence → lead to better calibration • Can improve generalization in ambiguous regions
	 • Forces model to treat all disagreements as certain • Overconfidence / poor calibration • Can overfit to arbitrary consensus → implicitly assume certainty in one GT	• Loss functions compatible with continuous labels • Slower convergence if uncertainty high • Choice of smoothing parameters matters
Validation	 • Easy evaluation metrics (Dice, IoU) • Easier interpretability by clinicians	• Can compare model probabilities directly to rater agreement • Better reflection of inherent ambiguity
	 • Metrics hide uncertainty (single number) • Apparent performance depends strongly on fusion method	• Harder to benchmark with standard metrics • Need thresholding or probabilistic metrics • May show lower peak Dice even if better calibrated
Best Use Case		
Low disagreement or when binary GT is required for deployment.		High disagreement or when uncertainty estimation / calibration is critical.

Outline

Motivation

Label fusion techniques


Method wrt Disagreement

Impact on model Training and Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Impact on Calibration



Fusion Choice



Training Target




Model Confidence



Calibration Quality

Impact on Calibration



Fusion Choice

Hard Labels:

- Treat all voxels as certain
- Model tends to become **overconfident in ambiguous areas**
- Calibration curves show underestimation of uncertainty

Soft Labels:

- Encode rater disagreement in the target
- Encourage model to **predict probabilities matching rater agreement** → real uncertainty
- Lead to improved Expected Calibration Error

Label fusion isn't just preprocessing — it defines the **(probability) target** your model learns, which directly **shapes calibration**.

Outline

Motivation

Label fusion techniques

Method wrt Disagreement

Soft vs Hard on Training/Validation

Impact on Calibration

Label Fusion Choice & Takeaways

Label Fusion Choice

	Why It Matters	Examples
Level of Disagreement	High disagreement → need richer models (STAPLE, soft labels) Low disagreement → simple consensus often sufficient.	Tumor margins vs. well-defined organs
Computational Resources & Time	Some methods (STAPLE, Bayesian models) require EM iterations or MCMC, can be slow for 3D volumes. Simple majority/union is fast .	Large CT volumes with many raters
Goal of the Model	If the model's main output is binary masks for deployment, simple fusion may suffice. If you need calibrated probabilities/uncertainty , prefer soft/probabilistic methods.	Clinical decision support vs. research prototype or other research fields
Complexity of the Problem	More complex anatomical regions or subtle findings → benefit from probabilistic/soft approaches . Well-defined , high-contrast structures → simpler methods .	Pancreatic tumor vs. lung parenchyma segmentation
Data Size & Number of Raters	Fewer raters → probabilistic methods may be less stable . More raters → benefits grow .	Small expert panel vs. crowdsourced annotations

Label Fusion Choice & Takeaways

- **Interrater variability is a valuable signal**, not just noise — it reflects real uncertainty in medical imaging. This uncertainty should be reflected in the predictions as well.
- Different **label fusion** methods offer **trade-offs** in **simplicity**, **computational cost**, and ability to **model uncertainty**.
- **Fusion choices** directly **shape** model training **targets** and strongly influence calibration, confidence, and clinical reliability.
- Choosing the right fusion strategy depends on **disagreement level**, **computational resources**, and the **intended goal** of the model — and this variability should be addressed not only during training but **also in evaluation**.
 - See spotlight presentation.

Thank you

Meritxell Riera i Marín, PhD candidate
m.riera@sycatechnologies.com