Part IV - B

Model calibration

Conformal Prediction

Adrian Galdran
RyC/ikerbasque Research Fellow
Tecnalia - Derio, Spain
adrian.galdran@tecnalia.com

Conformal Prediction - Contents

1. Motivation
2. Conformal Prediction: Ingredients
3. Conformal Predictions: Algorithm
4. Hands-On

1. Motivation

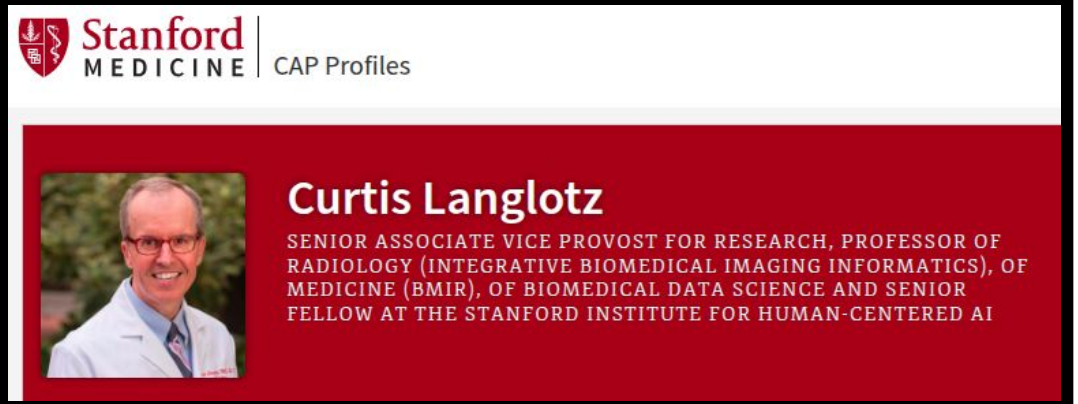

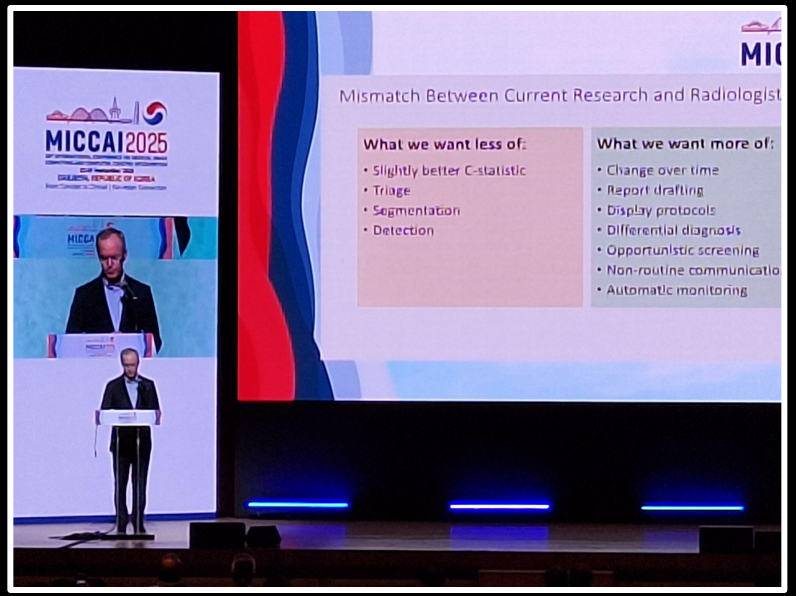

1. Motivation

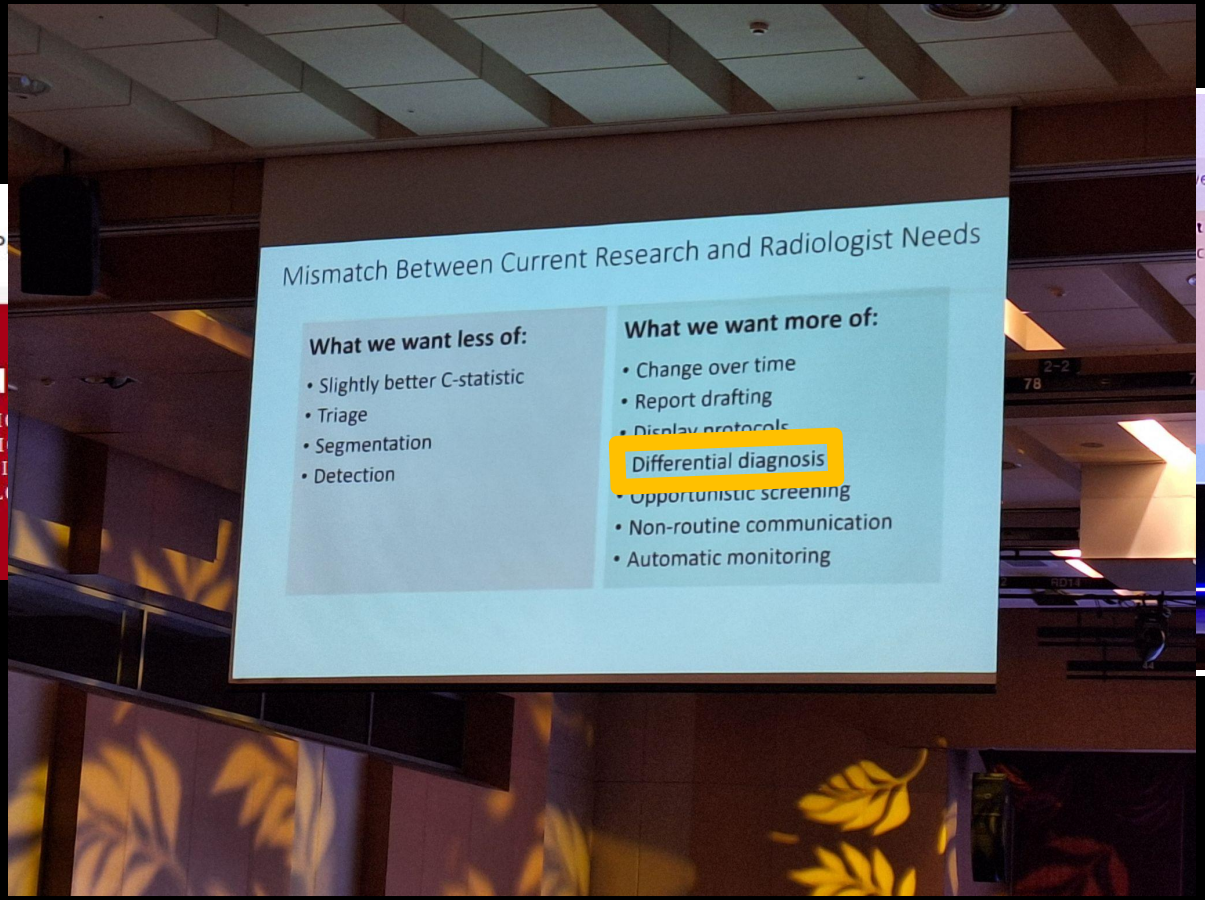

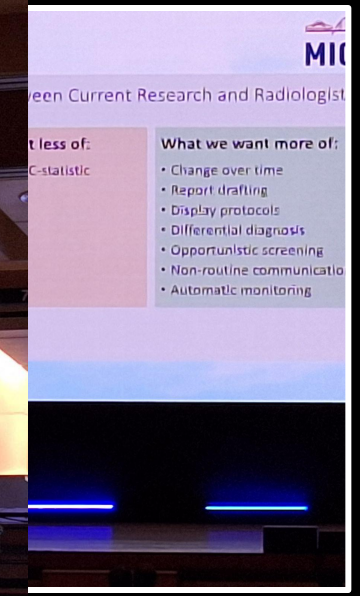

MICCAI2025
Daejeon
REPUBLIC OF KOREA

1. Motivation

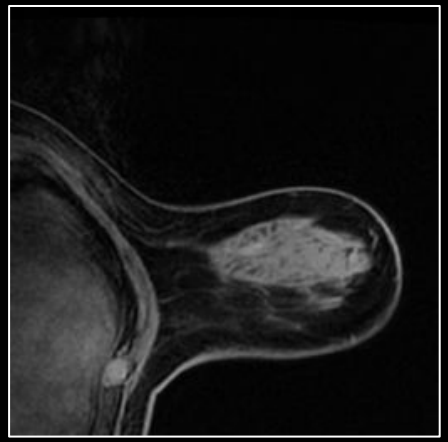

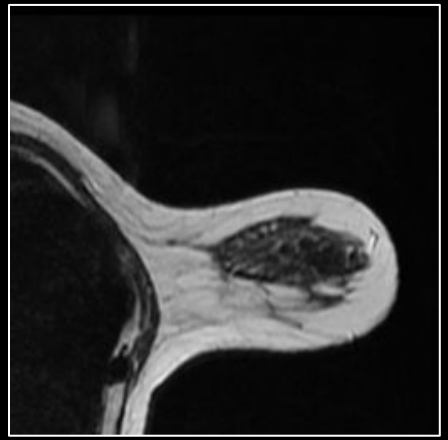

Dietzel M, Baltzer PAT. How to use the Kaiser score as a clinical decision rule for diagnosis in multiparametric breast MRI: a pictorial essay. Insights Imaging. 2018

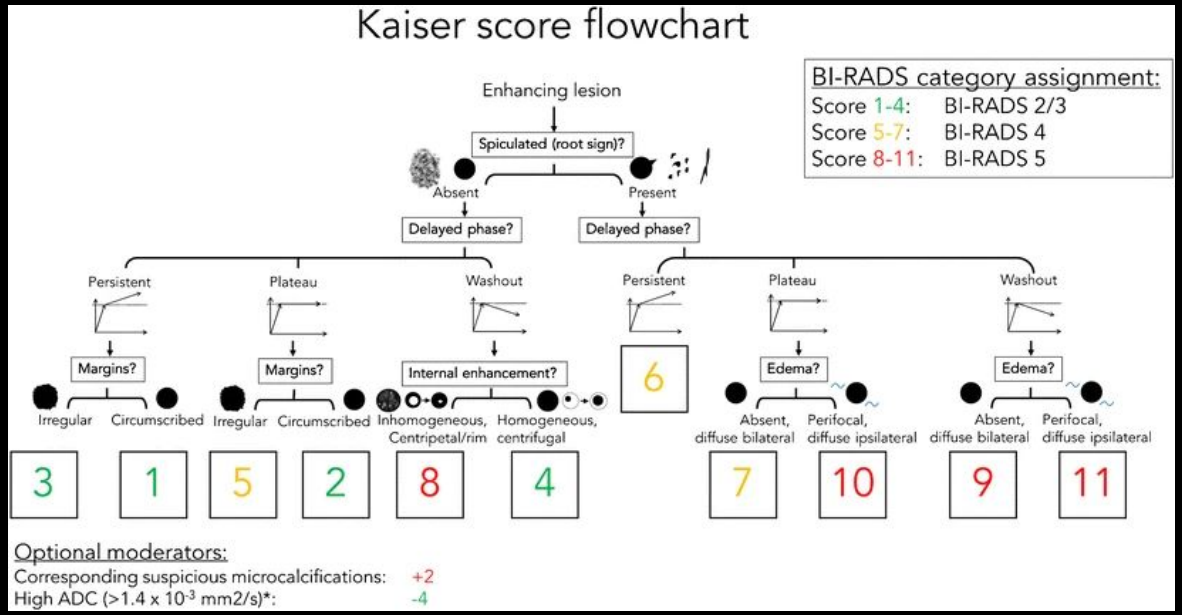

1. Motivation

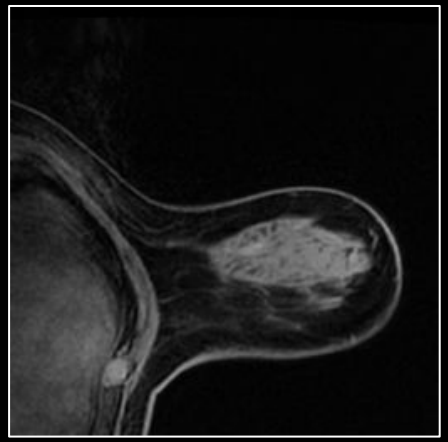

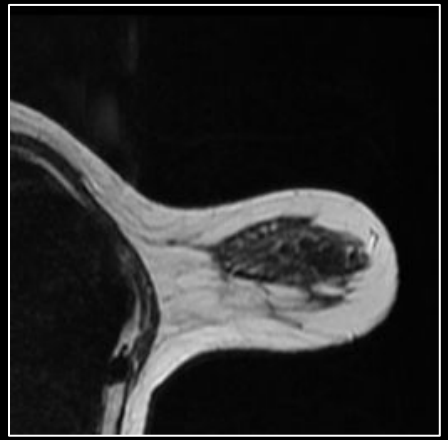

Dietzel M, Baltzer PAT. How to use the Kaiser score as a clinical decision rule for diagnosis in multiparametric breast MRI: a pictorial essay. Insights Imaging. 2018

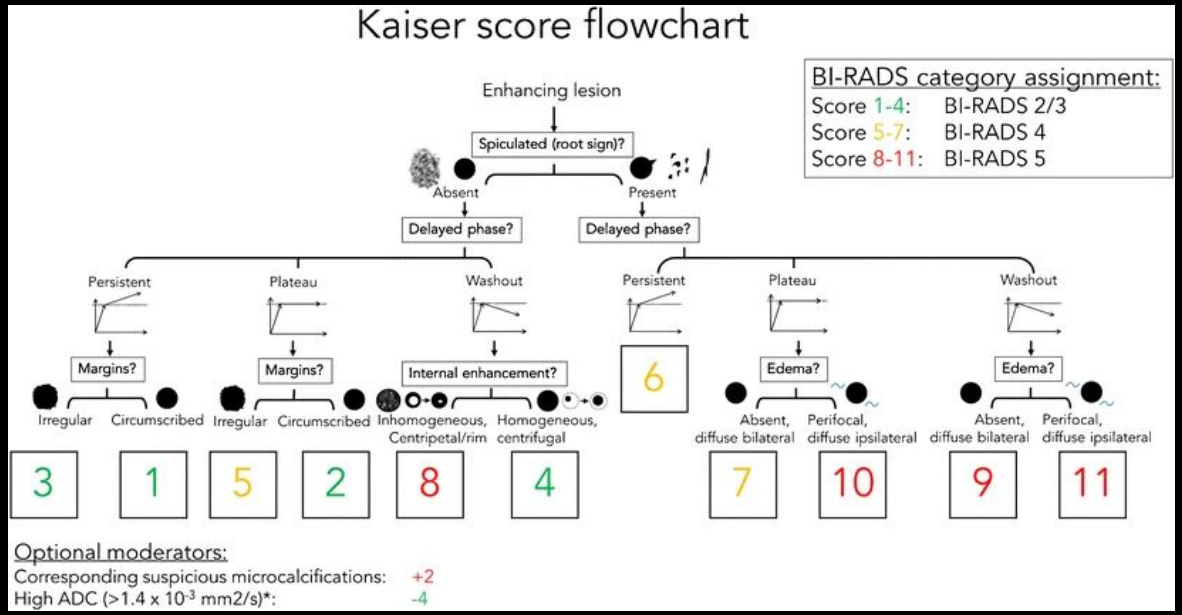

$$\mathbb{P}(\text{BI-RADS} \in \{7, 9, 10, 11\}) \geq 90\%$$

1. Motivation: Differential Diagnosis

Step 1: Do we see an enhancing lesion?

- On MRI, when you inject contrast, suspicious lesions often "light up" because they have abnormal blood vessels.
- So the first question is: *does the lesion actually enhance?* If yes → move on.

Step 2: Is there a spiculated margin ("root sign")?

- **Spiculation** means the lesion has spikes or radiating lines extending into the surrounding tissue — like roots of a tree.
- This is a strong red flag for cancer. If it's present, you follow the *right-hand branch*.
- If absent (the lesion looks smooth/rounded), you go to the *left-hand branch*.

Step 3: What happens in the delayed phase (contrast wash-out curve)?

This is the part you asked about — it's the **time course** of contrast enhancement.

- After contrast injection, radiologists watch how bright the lesion gets over time.
- There are three typical "curves":
  - **Persistent:** keeps getting brighter and brighter with time.
    → Usually benign (think of a sponge slowly soaking water).
  - **Plateau:** gets bright quickly, then levels off.
    → Suspicious (like tissue that soaks fast but then "caps out").
  - **Washout:** gets bright early, but then *fades* as contrast drains away.
    → Very suspicious for cancer (because malignant tumors often have "leaky" vessels).

So the **"delayed phase" check** is: which of these three time-curves does the lesion follow?

Step 4: If persistent or plateau → check margins

- **Margins** = edge of the lesion.
- Smooth (circumscribed) edges → usually benign.
- Irregular/jagged edges → more worrisome.

Step 5: If washout → check internal enhancement pattern

- Inside the lesion, how does the contrast distribute?
- If it's patchy, rim-shaped, or irregular → higher suspicion.
- If it's uniform or "centrifugal" (from inside out), less worrisome.

Step 6: If spiculation was present (right side of the tree) → check delayed phase again, then edema

- If the lesion is spiculated and the curve is persistent/plateau/washout, you still branch down.
- In later branches, radiologists look for **edema** — swelling in the tissue around the lesion.
- Edema often shows up in malignancy, so its presence increases the suspicion.

Putting it together

At the bottom of the tree, you land on a **Kaiser score number (1–11)**.

- **Low scores (1–4)**: likely benign (BI-RADS 2/3).
- **Middle scores (5–7)**: indeterminate but suspicious (BI-RADS 4).
- **High scores (8–11)**: very suspicious, likely malignant (BI-RADS 5).

2. Conformal Prediction: Ingredients

Suppose you've got an FDA-approved diagnostic model

$$\hat{\mathcal{M}}_y(\boldsymbol{x}) \sim \mathbb{P}(\boldsymbol{Y} = \boldsymbol{y} \,|\, \boldsymbol{X} = \boldsymbol{x}) \qquad \boldsymbol{y} \in \{1, \ldots, \boldsymbol{K}\} = \boldsymbol{\mathcal{Y}}$$

2. Conformal Prediction: Ingredients

Suppose you've got an FDA-approved diagnostic model

$$\hat{\mathcal{M}}_y(\boldsymbol{x}) \sim \mathbb{P}(\boldsymbol{Y} = \boldsymbol{y} \,|\, \boldsymbol{X} = \boldsymbol{x}) \qquad \boldsymbol{y} \in \{1, \ldots, \boldsymbol{K}\} = \mathcal{Y}$$

$$\hat{\mathcal{M}}_y(\boldsymbol{x}^*) = 7$$

2. Conformal Prediction: Ingredients

Suppose you've got an FDA-approved diagnostic model

$$\hat{\mathcal{M}}_y(\boldsymbol{x}) \sim \mathbb{P}(\boldsymbol{Y} = \boldsymbol{y} \,|\, \boldsymbol{X} = \boldsymbol{x}) \qquad \boldsymbol{y} \in \{1, \ldots, \boldsymbol{K}\} = \mathcal{Y}$$

$$\hat{\mathcal{M}}_y(\boldsymbol{x}^*) = 7$$

$$\hat{\mathcal{M}}_y(\boldsymbol{x}^*) \in \{7, 8, 9\}$$

$$\textit{with } p = 0.9$$

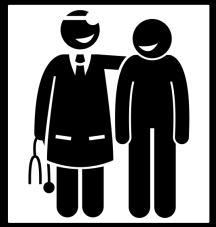

2. Conformal Prediction: Ingredients

Suppose you've got an FDA-approved diagnostic model

$$\hat{\mathcal{M}}_y(\boldsymbol{x}) \sim \mathbb{P}(\boldsymbol{Y} = \boldsymbol{y} \,|\, \boldsymbol{X} = \boldsymbol{x}) \qquad \boldsymbol{y} \in \{1, \ldots, \boldsymbol{K}\} = \mathcal{Y}$$

$$\hat{\mathcal{M}}_y(\boldsymbol{x}^*) = 7$$

$$\hat{\mathcal{M}}_y(\boldsymbol{x}^*) \in \{7, 8, 9\}$$

$$\textit{with } p = 0.9$$

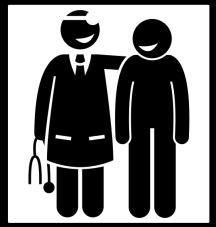

Predict a set $\mathcal{T}_{x^*} \subseteq \mathcal{Y}$ which contains y^* with high p

$$\textbf{Coverage}: \quad \mathbb{P}(y^* \in \mathcal{T}_{x^*}) \geq 1 - \alpha$$

MICCAI2025 Daejeon REPUBLIC OF KOREA

2. Conformal Prediction: Ingredients

Prediction Sets : $x^* \mapsto \mathcal{T}_{x^*} \subseteq \mathcal{Y} = \{1, \dots, \mathcal{K}\}$

Coverage : $\mathbb{P}(y^* \in \mathcal{T}_{x^*}) \geq 1 - \alpha$

---

2. Conformal Prediction: Ingredients

Prediction Sets : $x^* \mapsto \mathcal{T}_{x^*} \subseteq \mathcal{Y} = \{1, \ldots, K\}$

Coverage : $\mathbb{P}(y^* \in \mathcal{T}_{x^*}) \geq 1 - \alpha$

{ fox squirrel 0.99 }

{ fox squirrel 0.82, gray fox, 0.03, bucket, 0.02, rain barrel 0.02 }

{ marmot 0.30, fox squirrel 0.22, mink 0.18, weasel 0.16, beaver 0.03, polecat 0.01 }

Angelopoulos & Bates. A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification, arXiv:2107.07511, 2021

2. Conformal Prediction: Ingredients

Prediction Sets : $x^* \mapsto \mathcal{T}_{x^*} \subseteq \mathcal{Y} = \{1, \ldots, K\}$

Coverage : $\mathbb{P}(y^* \in \mathcal{T}_{x^*}) \geq 1 - \alpha$ **Efficient** sets

{ fox squirrel 0.99 }

{ fox squirrel 0.82, gray fox, 0.03, bucket, 0.02, rain barrel 0.02 }

{ marmot 0.30, fox squirrel 0.22, mink 0.18, weasel 0.16, beaver 0.03, polecat 0.01 }

Angelopoulos & Bates. A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification, arXiv:2107.07511, 2021

3. Conformal Predictions: Algorithm

Let's get back to our FDA-approved diagnostic model $\hat{\mathcal{M}}$

3. Conformal Predictions: Algorithm

Let's get back to our **FDA-approved** diagnostic model $\hat{\mathcal{M}}$

No retraining allowed

3. Conformal Predictions: Algorithm

Let's get back to our **FDA-approved** diagnostic model $\hat{\mathcal{M}}$

No retraining allowed, but you have some fresh data available:

$$(\texttt{bMRI}_1, y_1), \ldots, (\texttt{bMRI}_N, y_i) = \{(x_i, y_i)\}_{i=1}^{N} \sim \mathbb{P} \; i.i.d.$$

We call this **Calibration Set** (sorry about that)

3. Conformal Predictions: Algorithm

Let's get back to our **FDA-approved** diagnostic model $\hat{\mathcal{M}}$

No retraining allowed, but you have some fresh data available:

$$(\texttt{bMRI}_1, y_1), \ldots, (\texttt{bMRI}_N, y_i) = \{(x_i, y_i)\}_{i=1}^{N} \sim \mathbb{P} \ \textit{i.i.d.}$$

We call this **Calibration Set** (sorry about that)

How do we build these Conformal Prediction Sets?

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes

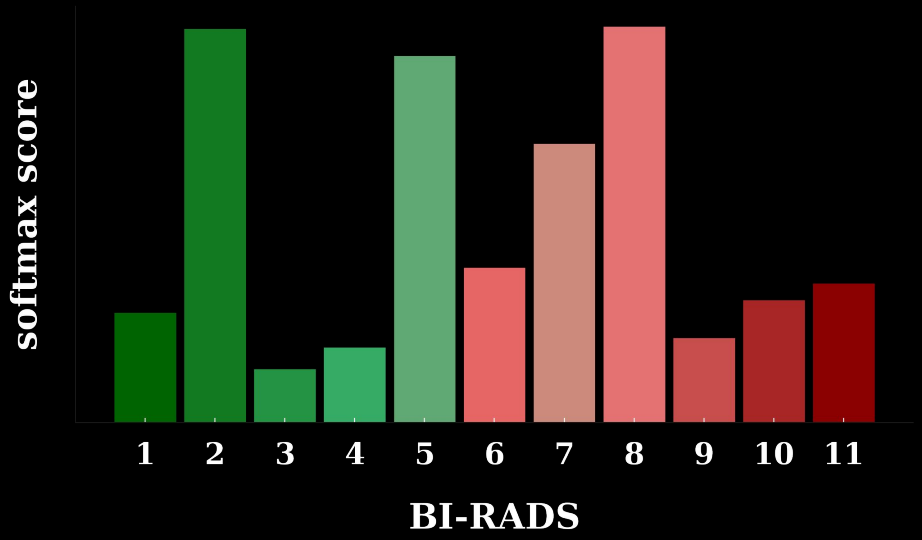

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, \}$

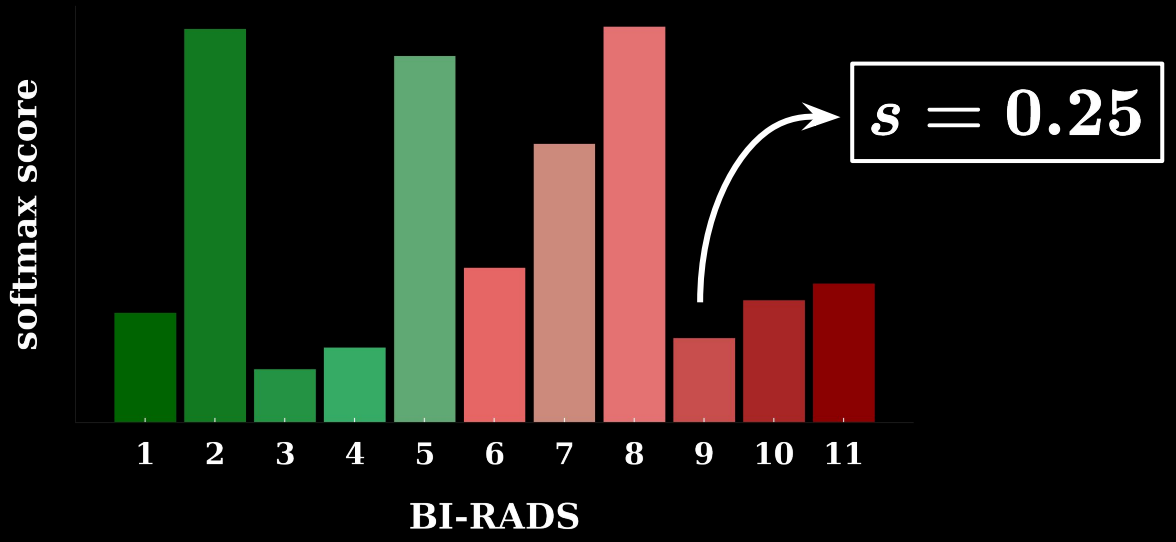

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, 0.90, \}$

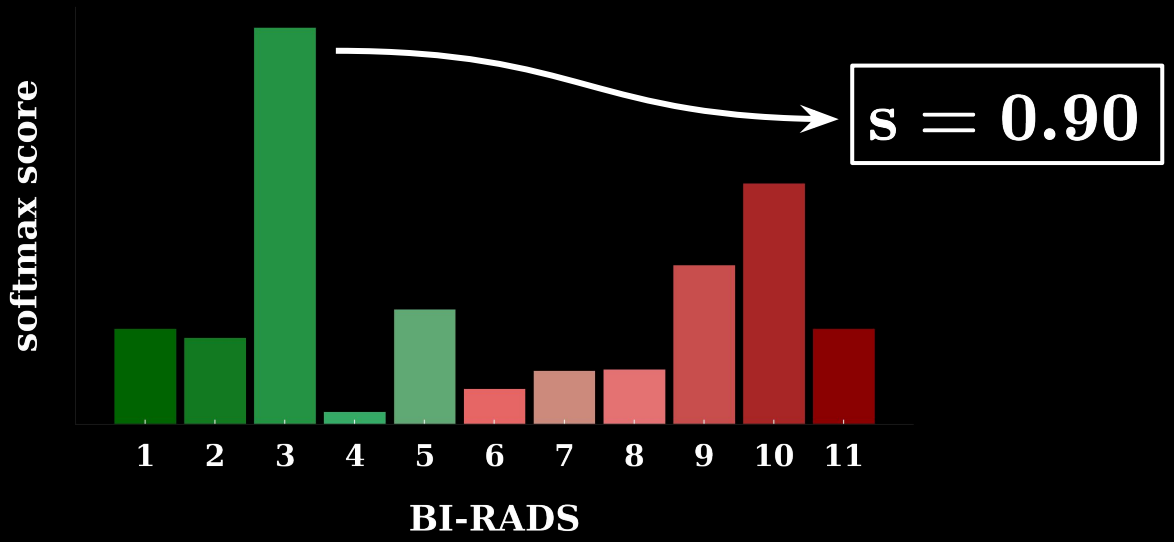

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, 0.90, \ldots, 0.65\}$

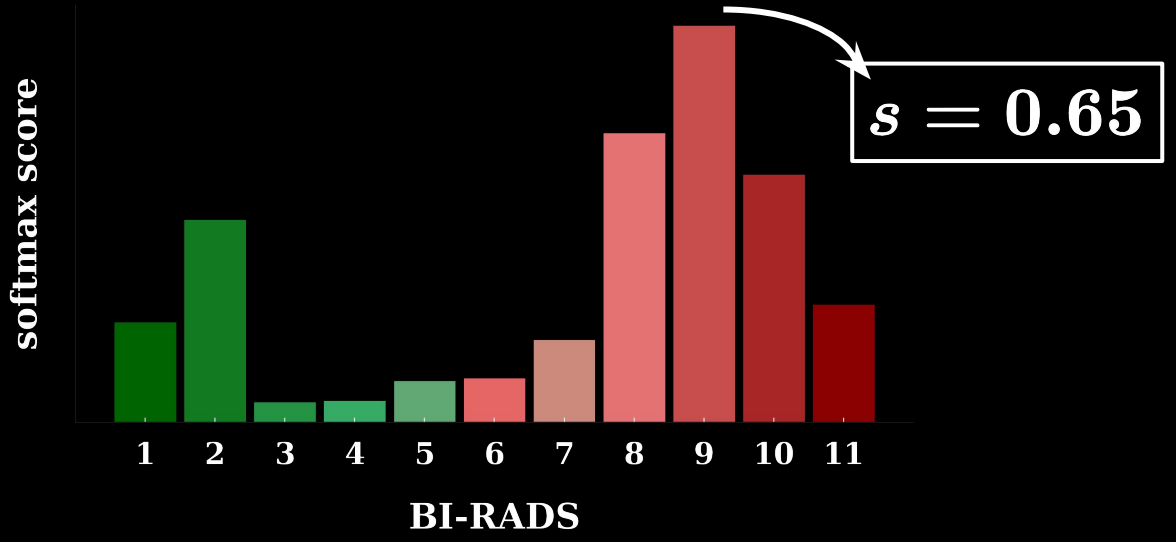

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, 0.90, \ldots, 0.65\}$
2. For a desired coverage of α, find a value $\hat{q_\alpha}$ such that you keep $1\text{-}\alpha$ of the scores in $\mathbb{E}$:

$$\hat{q_\alpha} = \texttt{np.quantile}([\mathtt{E}_1, \mathtt{E}_2, \ldots, \mathtt{E}_N], \alpha)$$

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, 0.90, \ldots, 0.65\}$
2. For a desired coverage of α, find a value $\hat{q}_\alpha$ such that you keep $1\text{-}\alpha$ of the scores in $\mathbb{E}$:

$$\hat{q}_\alpha = \texttt{np.quantile}([\mathrm{E}_1, \mathrm{E}_2, \ldots, \mathrm{E}_N], \alpha)$$

What happens if we use $\hat{q}_\alpha$ to build prediction sets in the calibration dataset?

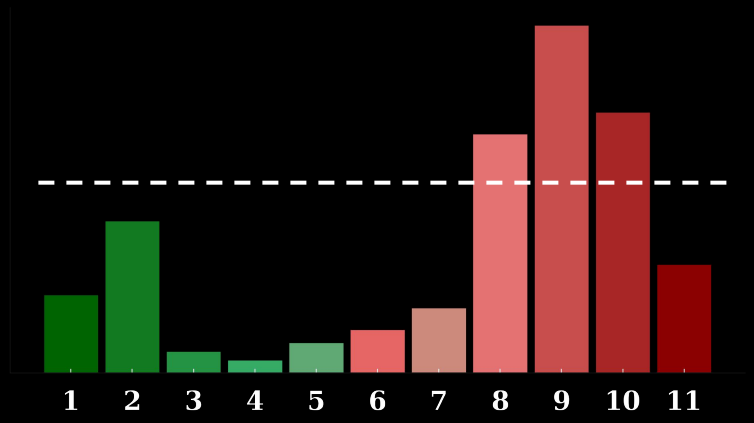

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25, 0.90, \ldots, 0.65\}$

2. For a desired coverage of α, find a value $\hat{q_\alpha}$ such that you keep $1\text{-}\alpha$ of the scores in $\mathbb{E}$:

$$\hat{q_\alpha} = \texttt{np.quantile}([\mathrm{E}_1, \mathrm{E}_2, \ldots, \mathrm{E}_N], \alpha)$$

3. If we use $\hat{q_\alpha}$ to build prediction sets on test data, we have <u>theoretically guaranteed coverage</u>

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25,\ 0.90,\ \ldots,\ 0.65\}$

2. For a desired coverage of α, find a value $\hat{q_\alpha}$ such that you keep $1\text{-}\alpha$ of the scores in $\mathbb{E}$:

$$\hat{q_\alpha} = \texttt{np.quantile}\left(\left[\mathrm{E}_1, \mathrm{E}_2, \ldots, \mathrm{E}_\mathrm{N}\right],\ \alpha\right)$$

3. If we use $\hat{q_\alpha}$ to build prediction sets on test data, we have <u>theoretically guaranteed coverage</u>, if data is <u>interchangeable</u>.

$$1\text{-}\alpha \leq \mathbb{P}\left(\mathbf{y}^* \in \mathcal{T}_{x^*}^{\hat{q_\alpha}}\right) \leq (1\text{-}\alpha) + \frac{1}{\mathrm{N}+1}$$

3. Conformal Predictions: Algorithm

1. Collect scores of correct classes $\mathbf{E}=\{0.25,\ 0.90,\ \ldots,\ 0.65\}$

2. For a desired coverage of α, find a value $\hat{q_\alpha}$ such that you keep $1\text{-}\alpha$ of the scores in $\mathbb{E}$:

$$\hat{q_\alpha} = \texttt{np.quantile}\left(\left[\mathrm{E}_1, \mathrm{E}_2, \ldots, \mathrm{E}_\mathrm{N}\right],\ \alpha\right)$$

3. If we use $\hat{q_\alpha}$ to build prediction sets on test data, we have <u>theoretically guaranteed coverage</u>. If data is <u>interchangeable</u>:

$$1\text{-}\alpha \leq \mathbb{P}\left(\mathbf{y}^* \in \mathcal{T}_{\boldsymbol{x}^*}^{\hat{q_\alpha}}\right) \leq (1\text{-}\alpha) + \frac{1}{\mathrm{N}+1}$$

MICCAI2025 Daejeon REPUBLIC OF KOREA

3.5 Beyond Coverage

Marginal Coverage of 90%

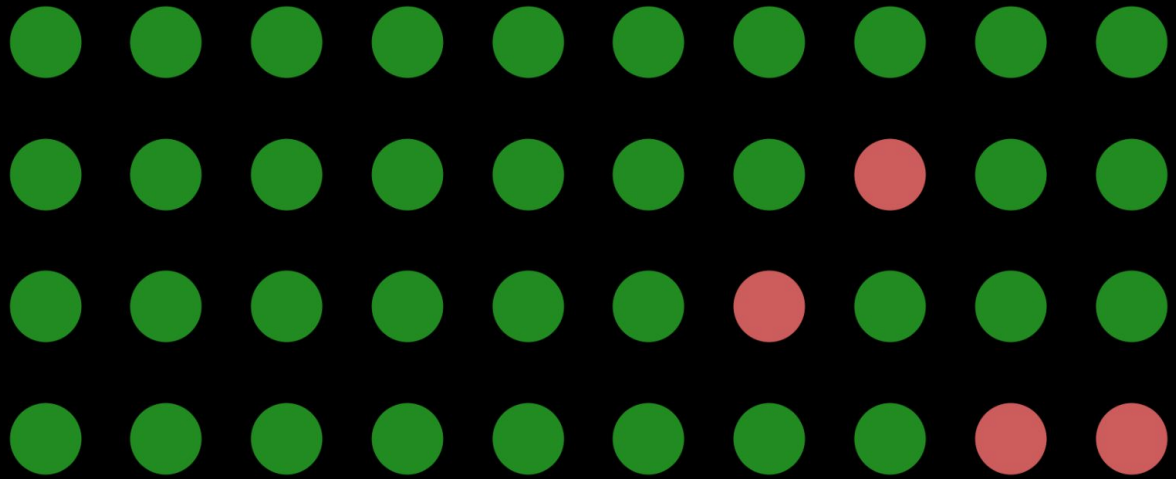

3.5 Beyond Coverage

Marginal Coverage of 90% **No Conditional Coverage**

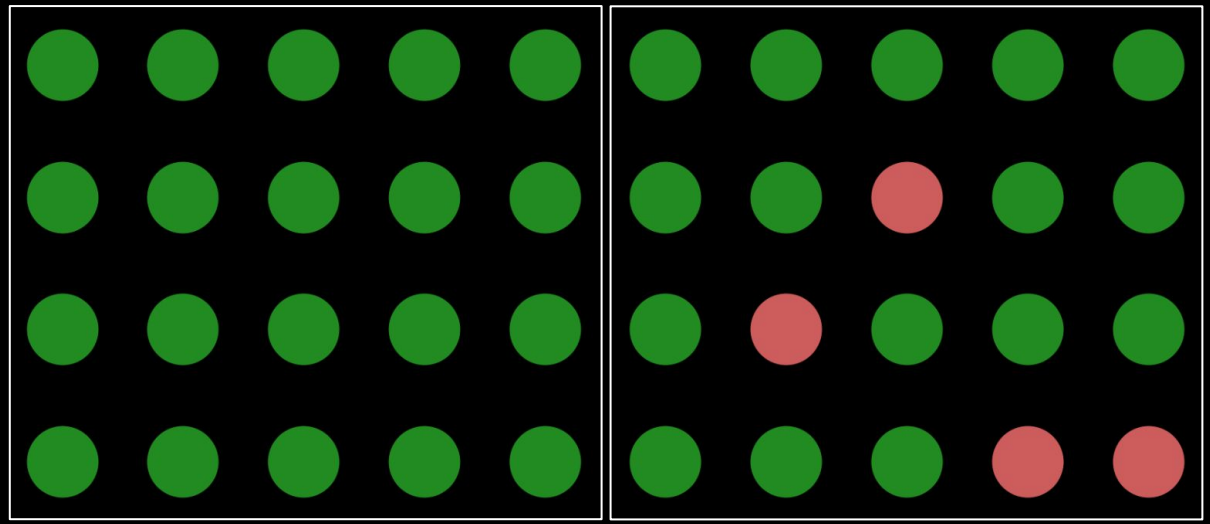

3.5 Beyond Coverage

Marginal Coverage of 90% **No Conditional Coverage**

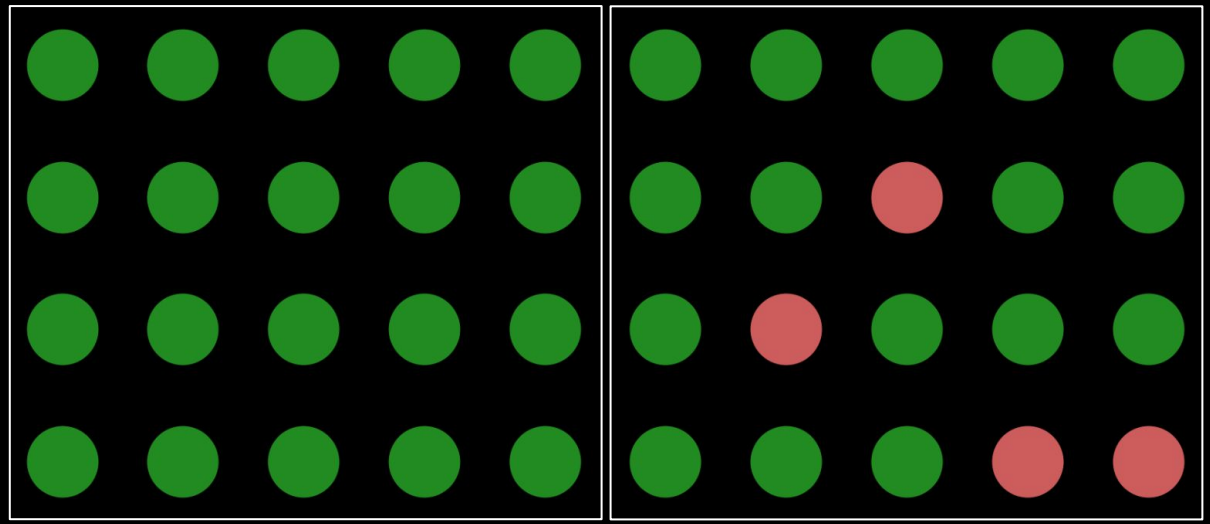

Check out also Conformal Risk Control

MICCAI2025 Daejeon REPUBLIC OF KOREA

4. Hands-On

Github repository:

https://github.com/agaldran/uqinmia-miccai

MICCAI2025 Daejeon REPUBLIC OF KOREA